

GAUSSIAN PROCESS REGRESSION FOR OUT-OF-SAMPLE EXTENSION

Oren Barkan, Jonathan Weill and Amir Averbuch

Tel Aviv University

ABSTRACT

Manifold learning methods are useful for high dimensional data analysis. Many of the existing methods produce a low dimensional representation that attempts to describe the intrinsic geometric structure of the original data. Typically, this process is computationally expensive and the produced embedding is limited to the training data. In many real life scenarios, the ability to produce embedding of unseen samples is essential. In this paper we propose a Bayesian non-parametric approach for out-of-sample extension. The method is based on Gaussian Process Regression and independent of the manifold learning algorithm. Additionally, the method naturally provides a measure for the degree of abnormality for a newly arrived data point that did not participate in the training process. We derive the mathematical connection between the proposed method and the Nystrom extension and show that the latter is a special case of the former. We present extensive experimental results that demonstrate the performance of the proposed method and compare it to other existing out-of-sample extension methods.

1. INTRODUCTION AND RELATED WORK

Dimensionality reduction methods are widely used in the machine learning community for high dimensional data analysis. Manifold learning is a subclass of non-linear dimensionality reduction algorithms. These algorithms attempt to discover the low dimensional manifold that the data points have been sampled from [1]. Many manifold learning algorithms produce an embedding of high dimensional data points in a low dimensional space. In this space, the Euclidean distance indicates the affinity between the original data points with respect to the manifold geometric structure. Typically, the embedding is produced only for the training data points, with no extension for out-of-sample points. Moreover, the process of computing the embedding usually involves expensive computational operations such as Singular Values Decomposition (SVD).

As a result, the application of manifold learning algorithms to massive datasets or data which is accumulated over time becomes impractical. Therefore, the out-of-sample extension (OOSE) problem is a major concern for manifold learning algorithms and over the years many methods have been proposed to alleviate this problem: Bengio et al. [2] proposed extensions for several well-known manifold learning algorithms. The extensions are based on the Nystrom extension [9], which has been widely used for manifold learning algorithms. In [3] the authors proposed to use the Nystrom extension of eigenfunctions of the kernel, however in order to maintain numerical stability, they used only the significant eigenvalues. As a result, the method might suffer from inconsistencies with the in-sample data. Bermanis et al. [4] suggested to alleviate the aforementioned problem by introducing a method for extending functions using a coarse-to-fine hierarchy of the multiscale decomposition of a Gaussian kernel. Recently, Aizenbud et al. [5] suggested an extension which is based on local Principal Component Analysis (PCA).

Fernandez et al. [6] proposed an extension of Laplacian Pyramids model that incorporates a modified Leave One Out Cross Validation (LOOCV), but avoids the large computational cost of the standard one. In [7], the authors proposed to extend the embedding to unseen samples by finding a rotation and scale transformations of the sample's nearest neighbors. Then, the embedding is computed by the applying these transformations to the unseen samples. Yang et al. [20] introduced a manifold learning technique that enables OOSE using regularization.

In this paper, we propose a general framework for OOSE which is based on Gaussian Process Regression (GPR) [8]. The method is independent of the manifold learning algorithm and provides a measure of abnormality for a given test instance with respect to the training instances. The outline of the method is as follows: Given a training data and a manifold learning algorithm, we first apply the algorithm on the training data and compute the corresponding embeddings. Then, we learn the hyperparameters for a

GPR model using the training data and the embeddings. Finally, given an unseen test instance and the trained GPR model, we produce a predictive distribution and set the embedding value to the mode of the distribution. Furthermore, the variance of the predictive distribution quantifies the degree of abnormality in the test instance. We analyze the mathematical connection between the proposed method and the Nystrom extension [9] and show that the latter is a special case of the former.

In the context of Bayesian statistics, Lawrence et al. [21] showed how Gaussian Process Latent Variable models can be generalized through back-constraints (GPLVMBC) to preserve local geometries. While our method learns to extend any given mapping, GPLVMBC is not designed to extend it, but produces a new mapping, which is different from the original one.

We evaluate the proposed method on several well-known manifold learning algorithms and both synthetic and real world datasets. We demonstrate its performance and show it manages to achieve competitive results when compared with other OOSE methods. The rest of the paper is organized as follows: Section 2 overviews Gaussian Processes and GPR. Section 3 describes the proposed method and discusses its relation to the Nystrom extension [9]. In Section 4 we present experimental results and in Section 5 we conclude.

2. GAUSSIAN PROCESS REGRESSION

Given a training set

$$D = \{(\mathbf{x}_i, y_i) \mid \mathbf{x}_i \in \mathbb{R}^N, y_i \in \mathbb{R}, i = 1, \dots, m\},$$

which consists of pairs of input vector and noisy predictions (\mathbf{x}_i, y_i) , Bayesian regression deals with computing a predictive distribution of y_* for a new test instance \mathbf{x}_* . Typically, the noise is assumed to be additive, independent and Gaussian such that the relation between the input and the output is given by

$$y_i = f(\mathbf{x}_i) + \varepsilon_i, \quad \varepsilon_i \sim N(0, \sigma^2), \quad (1)$$

where f is a function that comes to model the noise free relation between \mathbf{x}_i and y_i and $N(a, b)$ stands for the normal distribution with a mean a and a variance b .

A Gaussian Process (GP) is a stochastic process such that any finite subcollection of random variables has a multivariate Gaussian distribution. Gaussian Process Regression (GPR) is a non-parametric Bayesian regression model that assumes prior distribution of the function values such that $p(\mathbf{f} \mid \mathbf{x}_{1:m}) = N(\mathbf{0}, K_{\mathbf{ff}})$ where $\mathbf{f} = [f_1, \dots, f_m]^T$ ($f_i = f(\mathbf{x}_i)$) is a vector whose entries are the function values (note that these function values are

treated as random variables). $K_{\mathbf{ff}} \in \mathbb{R}^{m \times m}$ is a covariance matrix whose entries are computed by a covariance function $[K_{\mathbf{ff}}]_{ij} = \text{cov}(f_i, f_j) = k(\mathbf{x}_i, \mathbf{x}_j)$. Then, for a given test vector \mathbf{x}_* the predictive distribution of y_* can be computed by marginalizing out the function values \mathbf{f}

$$p(f_* \mid \mathbf{y}) = \int p(f_*, \mathbf{f} \mid \mathbf{y}) d\mathbf{f} = p(\mathbf{y})^{-1} \int p(\mathbf{y} \mid \mathbf{f}) p(f_*, \mathbf{f}) d\mathbf{f} \quad (2)$$

where the last transition in Eq.(2) follows Bayes rule and the fact that \mathbf{y} is conditionally independent of f_* given \mathbf{f} . Since both factors in the last integral in Eq.(2) have the following Gaussian distributions

$$p(f_*, \mathbf{f}) = N\left(\mathbf{0}, \begin{bmatrix} K_{\mathbf{ff}} & K_{f_*\mathbf{f}} \\ K_{\mathbf{f}f_*} & K_{f_*f_*} \end{bmatrix}\right), \quad p(\mathbf{y} \mid \mathbf{f}) = N(\mathbf{f}, \sigma^2 \mathbf{I}),$$

a closed form expression exists [8] for the predictive distribution

$$p(f_* \mid \mathbf{y}) = N(\mu_*, \sigma_*^2), \quad \mu_* = K_{f_*\mathbf{f}} \mathbf{A} \mathbf{y}, \quad (3)$$

$$\sigma_*^2 = K_{f_*f_*} - K_{f_*\mathbf{f}} \mathbf{A} K_{\mathbf{f}f_*}, \quad \mathbf{A} = (K_{\mathbf{ff}} + \sigma^2 \mathbf{I})^{-1}.$$

Therefore, training a GPR model amounts to the computation of \mathbf{A} and $\mathbf{A} \mathbf{y}$. The computational complexity of the training procedure is dominated by a matrix inversion which is $O(n^3)$. Then, the prediction for a new test instance \mathbf{x}_* is given by the mode of $p(f_* \mid \mathbf{y})$, which is the mean μ_* in the case of the Gaussian distribution. The variance σ_*^2 serves as a measure of uncertainty in the prediction.

3. GPR BASED OOSE

Given a manifold learning algorithm M and a training set $\mathbf{X} = \{\mathbf{x}_i\}_{i=1}^m \subset \mathbb{R}^N$, we apply M to \mathbf{X} and compute the corresponding low dimensional embedding $\mathbf{Y} = \{\mathbf{y}_i\}_{i=1}^m \subset \mathbb{R}^d$ ($d \ll n$). Then, for each dimension $1 \leq j \leq d$, independently, we form a new training set $D_j = \{(\mathbf{x}_i, y_{ij}) \mid \mathbf{x}_i \in \mathbb{R}^N, y_{ij} \in \mathbb{R}, i = 1, \dots, m\}$ and train a separate GPR model. Then, given an unseen test example \mathbf{x}_* , we predict (Eq. (3)) its embedding and the measure of uncertainty in the predictions by $\mathbf{y}_* = \boldsymbol{\mu}_* = [\mu_{*1}, \dots, \mu_{*d}]^T$ and $\boldsymbol{\sigma}_*^2 = [\sigma_{*1}^2, \dots, \sigma_{*d}^2]^T$, respectively. As the variance increases, our confidence in the prediction decreases and \mathbf{x}_* might be considered as anomaly with respect to training set \mathbf{X} .

In this work we use the squared exponential covariance

function (kernel) $k(\mathbf{x}_i, \mathbf{x}_j) = \exp(-\tau^{-2} \|\mathbf{x}_i - \mathbf{x}_j\|_2^2)$, where τ is a hyperparameter that determines the width of the kernel (in our experiments, we evaluated several other kernels and they did not provide any significant improvement). An additional hyperparameter is the noise variance σ^2 in Eq. (1). The hyperparameters can be optimized with respect to D_j (note that the optimization is done for each GPR model j in separate). One option is to compute type II Maximum Likelihood (ML) estimates for the hyperparameters with respect to D_j . In the literature, this method is named as marginal likelihood [8]. Another approach is to apply cross validation. Fortunately, a close form expressions for LOOCV and its gradients exist [8] and the hyperparameters can be optimized with the Conjugate Gradient method. In this work, we use LOOCV for hyperparameter optimization. The main reason we chose this approach is that the marginal likelihood method is more prone to overfitting [8]. The algorithm is summarized in Fig.1.

3.1 The relation between GPR and the Nystrom extension

Many manifolds learning methods are cast in the same framework [2], where the computation of the embedding of the training data points is obtained by eigendecomposition of a (normalized) kernel matrix. Therefore, for a given training set $\mathbf{X} = \{\mathbf{x}_i\}_{i=1}^m \subset \mathbb{R}^N$ the kernel matrix is computed by $K_{ij} = k(\mathbf{x}_i, \mathbf{x}_j)$ (this might be followed by a subsequent normalization). Then, the eigendecomposition of K is carried out to form the following relation

$$K = \mathbf{Y} \mathbf{\Lambda} \mathbf{Y}^T, \quad (4)$$

where $\mathbf{\Lambda}$ and \mathbf{Y} are a diagonal matrix with the n eigenvalues on its diagonal and their corresponding column eigenvectors, respectively. Note that K is a real symmetric matrix and hence $\mathbf{Y}^T = \mathbf{Y}^{-1}$. Finally, the embedding of \mathbf{x}_i is obtained by the i -th row of \mathbf{Y} . We can rewrite Eq.(4) as

$$\mathbf{Y}_{ij} = \lambda_j^{-1} K_{i\mathbf{x}} \mathbf{y}_j = \lambda_j^{-1} \sum_{z=1}^n k(\mathbf{x}_i, \mathbf{x}_z) \mathbf{Y}_{zj}, \quad (5)$$

where \mathbf{y}_j is the j -th column eigenvector in \mathbf{Y} and $K_{i\mathbf{x}}$ is the i -th row in K . In other words, the embedding for each data point in the training set is determined by a linear combination of the embeddings of all the other training data points multiplied by the inverse of the corresponding eigenvalue. The linear coefficients are the scaled kernel values. For the sake of simplicity, we limit the discussion to a single dimensional embedding,

GPR based OOSE

Training phase

Input:

M - manifold learning algorithm

$\mathbf{X} = \{\mathbf{x}_i\}_{i=1}^m$ - training set

d - target dimensionality

K - kernel function

Output:

$G = \{G_i\}_{i=1}^d$ - set of trained GPR models for each target dimension.

1. Compute the embedding $\mathbf{Y} = \{\mathbf{y}_i\}_{i=1}^m$ using M and \mathbf{X} .

2. For $j \leftarrow 1$ to d

2.1. $D_j \leftarrow \{(\mathbf{x}_i, y_{ij}) \mid i = 1, \dots, m\}$

2.2. Update $K^{(j)}$ and σ^2 using LOOCV [8].

2.3. $\mathbf{v} \leftarrow [y_{1j}, \dots, y_{mj}]^T$

2.4. $\mathbf{A}_j \leftarrow (K_{\mathbf{ff}}^{(j)} + \sigma_j^2 \mathbf{I})^{-1}$ (Eq. (3))

2.5. $\mathbf{w}_j \leftarrow \mathbf{A}_j \mathbf{v}$

2.6. $G_j \leftarrow \{\mathbf{A}_j, \mathbf{w}_j, K^{(j)}\}$

Test phase

Input:

\mathbf{x}_* - test instance.

$G = \{G_i\}_{i=1}^d$ - set of trained GPR models for each target dimension.

Output:

\mathbf{y}_* - the prediction for \mathbf{x}_*

σ_* - measure of uncertainty for each of \mathbf{x}_* entries

1. For $j \leftarrow 1$ to d

1.1. $\mathbf{y}_{*j} \leftarrow K_{f,\mathbf{f}}^{(j)} \mathbf{w}_j$

1.2. $\sigma_{*j} \leftarrow K_{f,f_*}^{(j)} - K_{f,\mathbf{f}}^{(j)} \mathbf{A}_j K_{\mathbf{f},f_*}^{(j)}$

Figure 1: GPR based OOSE algorithm

the generalization for multidimensional embedding is straightforward.

The Nystrom extension proposes to compute the embedding \mathbf{y}_* for a new test instance \mathbf{x}_* by

$$\mathbf{y}_{*j} = \lambda_j^{-1} K_{*\mathbf{x}} \mathbf{y}_j = \lambda_j^{-1} \sum_{z=1}^n k(\mathbf{x}_*, \mathbf{x}_z) \mathbf{Y}_{zj} \quad (6)$$

which amounts to the application of the kernel for each data point in the training set \mathbf{X} with respect to \mathbf{x}_* , followed by a dot product with \mathbf{y}_j .

Assuming a noise free GPR model with an identical kernel, the following relation holds: $y_i = K_{i\mathbf{x}} K^{-1} \mathbf{y}$ and the prediction in Eq.(3) reduces to

$$\mathbf{y}_{*j} = \mu_* = K_{*X} K^{-1} \mathbf{y}_j. \quad (7)$$

By using Eq. (4) we have

$$K^{-1} = (\mathbf{Y} \mathbf{A} \mathbf{Y}^T)^{-1} = \mathbf{Y} \mathbf{A}^{-1} \mathbf{Y}^T. \quad (8)$$

By combining Eqs. (7) and (8) we get

$$\mathbf{y}_{*j} = K_{*X} \mathbf{Y} \mathbf{A}^{-1} \mathbf{Y}^T \mathbf{y}_j = K_{*X} \mathbf{Y} \mathbf{A}^{-1} \mathbf{e}_j = \lambda_j^{-1} K_{*X} \mathbf{y}_j. \quad (9)$$

where the second transition is due to the fact that \mathbf{Y} 's columns are orthonormal and \mathbf{e}_j is the standard basis vector j . Notice that the predictions in Eqs. (6) and (9) are identical. Hence, the Nystrom extension is equivalent to a noise free GPR model with no hyperparamters optimization.

4. EXPERIMENTAL SETUP AND RESULTS

In this section we provide experimental results that demonstrate the performance of our proposed method and compare it to other existing OOSE methods.

4.1 The experimental workflow

The workflow of the experiments is as follows: Given a manifold learning algorithm M , OOSE method O and a dataset X , we apply M to X and derive corresponding embeddings Y . Then, we randomly divide (X, Y) to training and test sets $R = (X_{train}, Y_{train})$ and $Q = (X_{test}, Y_{test})$, respectively. The division is done according to a specific portion ρ (ρ is the fraction of data points assigned to R the rest are assigned to Q). Then, using O , R and X_{test} we produce embeddings \tilde{Y}_{test} . Finally, we measure the accuracy of the extension by the Root Mean Squared Error (RMSE) measure

$$RMSE(Y, \tilde{Y}) = \left(n^{-1} \sum_{i=1}^n \|y_i - \tilde{y}_i\|^2 \right)^{-1/2}.$$

We repeat the above procedure ten times (for different random divisions, R and Q) to produce a series of RMSE scores and determine the final RMSE as the series average. Note that our evaluation is similar to the other previous OOSE works [5-7], except for the fact we add the parameter ρ that challenges the evaluated methods with variable training set sizes. We will use the notations defined here throughout this section.

4.2 OOSE methods

We compare our proposed method to the Nystrom extension method and several OOSE methods that were recently developed and shown to overcome some of the limitations of the Nystrom extension. The methods are: Multiscale extension [4], Adaptive Laplacian Pyramids (ALP) [6] and

PCA based OOSE (POOS) [5]. All of the methods provide an OOSE scheme which is independent of the manifold learning algorithm.

4.3 Manifold learning algorithms

We evaluate the performance of the OOSE methods for several well-known manifold learning algorithms: Diffusion Maps (DM) [3], ISOMAP [11], Laplacian Eigenmaps (LE) [10], Local Linear Embeddings (LLE) [12] and Multidimensional Scaling (MDS) [13]. We set the methods hyperparameters according to the evaluated dataset: For all datasets we applied ISOMAP, LE, LLE and MDS with the same nearest neighbor value $k=8$. For the DM method, we adjust the neighborhood value according to the median of the squared Euclidean distances between the data points. For 3-dimensional datasets the target dimensionality for the manifold learning methods was set to 2. For high dimensional datasets the target dimension was chosen according to the spectral decay for each manifold learning algorithm, separately.

4.4 Datasets

We use various synthetic and real datasets: Swiss roll [11], Swiss hole [15], Corner planes [16], Punctured sphere [12], Twin peaks [12], 3D Clusters [16], Toroidal Helix [3], Faces [11], MNIST [17] and USPS [18]. Each dataset poses a different challenge for the manifold learning algorithm. All datasets were fed into the manifold learning algorithms and OOSE schemes as are, without any further preprocessing. From the MNIST and USPS datasets we randomly drew a collection of 2000 and 1000 images, respectively of the same digit. For each synthetic dataset we generated a set of 1000 data points.

4.5 Experiment 1

Our first experiment is designed to visualize the error obtained by a GPR based OOSE. To this end, we created a Swiss roll [11] with 1000 data points (Fig.(4) left bottom corner) and produced a corresponding embedding using each of the manifold learning algorithms, separately. Then, we created R and Q sets with $\rho=0.1$ as described in Section 4.1 and trained a GPR model for each of the manifold learning algorithms using R . Figure 2 presents the true embedding Y_{test} and the predictions \tilde{Y}_{test} that were produced for X_{test} using GPR. As we can see, even for a small value of ρ , the predictions managed to preserve a small amount of noise and follow the same structure of the true embeddings Y_{test} .

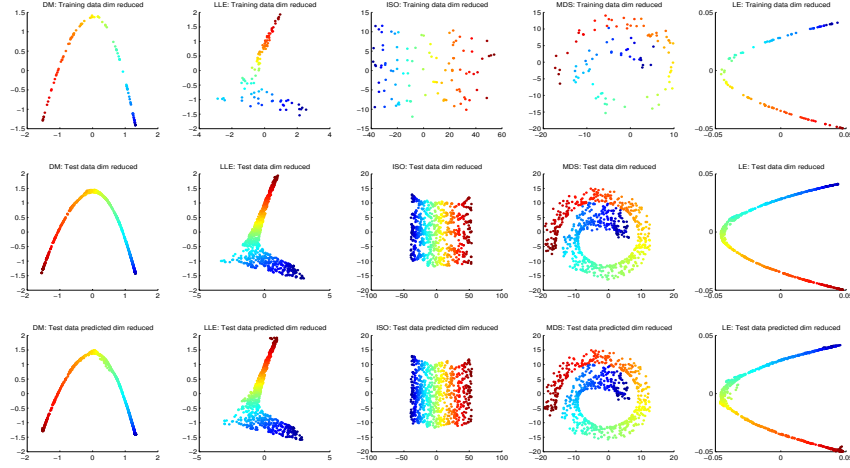


Figure 2: GPR based out-of-sample extension for various manifold learning methods and datasets. See Section 4.5 for details.

4.6 Experiment 2

This experiment is designed to evaluate the OOSE methods, each time on a specific pair of a manifold learning algorithm and a dataset. To this end, given a pair (M, X) we generate R and Q and compute average RMSE values for each O (see Section 4.1 for notations and further explanation). We repeat the experiment for increasing values of ρ starting from 0.05 to 0.8. Then for each O we plot a graph of the log RMSE as a function of ρ . We used the parameters that were specified in Section 4.3.

The results are presented in Fig. 3. (we did not add labels to the axis, since y values are measured relatively to the competitor methods, rather than their absolute values and axis x is ρ value which is clear from the context). Figure 3 is a graph table in which the (i, j) entry corresponds to a specific pair (M, X) . The pairs are clear from the row and columns labels. As we can see GPR produces the lowest RMSE values for most of the configurations followed by POOS as the second best method. The ALP seems to perform the worse, we conclude that it is due overfitting (in the ALP algorithm, a parameter is learnt from the training data and then used in test phase [6]). The reader might notice that some of the RMSE graphs are increasing in certain late intervals (mainly for real datasets), this might be explained by outliers or instability of the manifold learning algorithm: sometimes few points in the embedding are disconnected from the rest. Therefore, as ρ increases, the chance of these points to be included in R increases as well.

4.7 Experiment 3

We experimented with using our proposed model for an anomaly detection task on the DARPA Intrusion Detection Evaluation Data Set [19]. Each instance in this dataset has 14 features based on network traffic. Every instance is associated with standard network activity or a network attack and labeled accordingly. First, each dimension in the training set was mapped to $[0, 1]$ using a constant scale factor. The scale factors were saved to apply the same scaling to the test instances. Next, the training data was reduced to 2 dimensions using Diffusion Maps with a Gaussian kernel with a neighborhood parameter $\varepsilon = 0.7$ (which was computed using the median of the average $k=5$ nearest neighbor distances). To decide whether a test instance is an anomaly, we use our proposed OOSE method to obtain a distribution over the 2-dimensional diffusion space of a lower dimensional vector corresponding to the test instance. The final decision is made by comparing the variance of this Gaussian distribution to a threshold. The threshold was learned by holding out 20% of the training set and optimizing the prediction accuracy on the hold out set. Using this approach, we obtained an accuracy of over 99%.

5. CONCLUSION

In this paper, we proposed a Bayesian nonparametric method for OOSE. The method is based on GPR. We analyzed the relation between the Nystrom extension and GPR and showed that the former is a special case of the latter. We validated our proposed method in a series of experiments that demonstrated its performance and compared it to other OOSE methods.

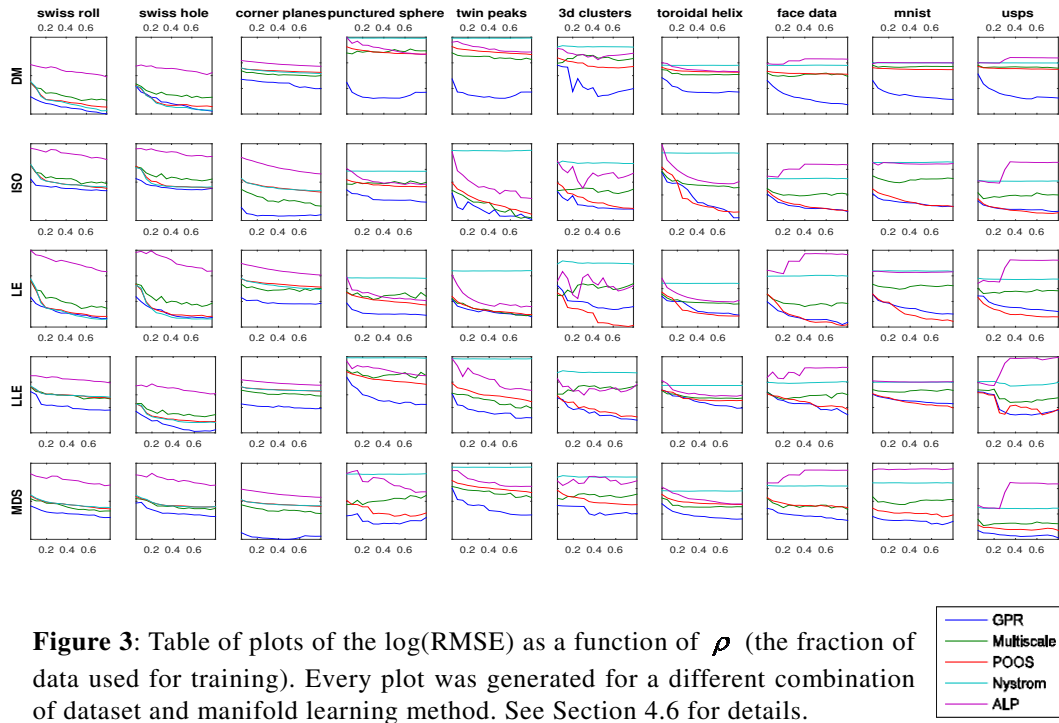


Figure 3: Table of plots of the log(RMSE) as a function of ρ (the fraction of data used for training). Every plot was generated for a different combination of dataset and manifold learning method. See Section 4.6 for details.

6. REFERENCES

- [1] van der Maaten, Laurens JP, Eric O. Postma, and H. Jaap van den Herik. "Dimensionality reduction: A comparative review." JMLR 10.1-41 (2009): 66-71.
- [2] Bengio, Yoshua, Jean-François Paiement, Pascal Vincent, Olivier Delalleau, Nicolas Le Roux, and Marie Ouimet. "Out-of-sample extensions for lle, isomap, mds, eigenmaps, and spectral clustering." In NIPS (2004): 177-184.
- [3] Lafon, Stéphane S. "Diffusion maps and geometric harmonics." PhD diss., Yale University, 2004.
- [4] Bermanis, Amit, Amir Averbuch, and Ronald R. Coifman. "Multiscale data sampling and function extension." Applied and Computational Harmonic Analysis 34, no. 1 (2013): 15-29.
- [5] Aizenbud, Yariv, Amit Bermanis, and Amir Averbuch. "PCA-Based Out-of-Sample Extension for Dimensionality Reduction.", preprint (2013).
- [6] Fernández, Ángela, Neta Rabin, Dalia Fishelov, and José R. Dorronsoro. "Auto-adaptative Laplacian Pyramids for High-dimensional Data Analysis." arXiv preprint arXiv:1311.6594 (2013).
- [7] Strange, Harry, and Reyer Zwiggelaar. "A generalised solution to the out-of-sample extension problem in manifold learning." In AAAI. 2011.
- [8] Rasmussen, Carl Edward. "Gaussian processes for machine learning." (2006).
- [9] Williams, Christopher, and Matthias Seeger. "Using the Nyström method to speed up kernel machines." In NIPS, pp. 682-688. 2001.
- [10] Belkin, Mikhail, and Partha Niyogi. "Laplacian eigenmaps and spectral techniques for embedding and clustering." In NIPS, vol. 14, pp. 585-591. 2001.
- [11] Tenenbaum, Joshua B., Vin De Silva, and John C. Langford. "A global geometric framework for nonlinear dimensionality reduction." Science 290, no. 5500 (2000).
- [12] Roweis, Sam T., and Lawrence K. Saul. "Nonlinear dimensionality reduction by locally linear embedding." Science 290, no. 5500 (2000): 2323-2326.
- [13] Cox, Trevor F., and Michael AA Cox. *Multidimensional scaling*. CRC Press, 2000.
- [14] Wilson, Andrew Gordon, and Ryan Prescott Adams. "Gaussian process kernels for pattern discovery and extrapolation." arXiv preprint arXiv:1302.4245 (2013).
- [15] Donoho, David L., and Carrie Grimes. "Hessian eigenmaps: Locally linear embedding techniques for high-dimensional data." Proceedings of the National Academy of Sciences 100, no. 10 (2003): 5591-5596.
- [16] <http://www.math.ucla.edu/~wittman/mani>
- [17] LeCun, Yann, Léon Bottou, Yoshua Bengio, and Patrick Haffner. "Gradient-based learning applied to document recognition." Proceedings of the IEEE 86, no. 11 (1998): 2278-2324.
- [18] Vapnik, Vladimir Naumovich, and Vladimir Vapnik. *Statistical learning theory*. Vol. 1. New York: Wiley, 1998.
- [19] Lippmann, Richard P. et al. "Evaluating intrusion detection systems: The 1998 DARPA off-line intrusion detection evaluation." In DARPA Information Survivability Conference and Exposition, 2000.
- [20] Yang, Yi, Feiping Nie, Shiming Xiang, Yueting Zhuang, and Wenhua Wang. "Local and Global Regressive Mapping for Manifold Learning with Out-of-Sample Extrapolation." In AAAI. 2010.
- [21] Lawrence, Neil D., and Joaquin Quiñero-Candela. "Local distance preservation in the GP-LVM through back constraints." In ICML, pp. 513-520. 2006.